

Lecture 1.2 : Strings

Introduction

- A string is simply a *sequence* of characters. The string type is a *collection type* meaning it contains a number of objects (characters in this case) that can be treated as a single object. The string type is a particular type of collection type called a *sequence type*. This means each constituent object (here character) in the collection occupies a specific numbered location within it i.e. elements are *ordered*.
- Python strings are typically enclosed in single or double quotes. (Pick a style and be consistent throughout and across your programs.)

```
name1 = "Jimmy Murphy"
name2 = 'Mary Kelly'
print(name1)
print(name2)
```

```
Jimmy Murphy
Mary Kelly
```

- What if you want to include quotes in your string? You have a couple of options: You can enclose your string in the other kind of quote or you can *escape* the quote with a backslash (nullifying its role as a string delimiter).

```
name3 = "Nora O'Neill"
name4 = 'Sally O\'Brien'
print(name3)
print(name4)
```

```
Nora O'Neill
Sally O'Brien
```

- Python strings can contain *non-printing* characters (such as a `\n` which causes a new line to be emitted when printing the string).

```
rhyme = "Humpty Dumpty sat on a wall,\nHumpty Dumpty had a great fall."
print(rhyme)
```

```
Humpty Dumpty sat on a wall,
Humpty Dumpty had a great fall.
```

String indexing

- As mentioned, a string is a *sequence* type. This means each member object (character in this case) occupies a numbered position in the collection and can be accessed by *indexing* the sequence at that index.
- For example, the characters of the string 'This is a sentence.' reside at the indices indicated here:

-19	-18	-17	-16	-15	-14	-13	-12	-11	-10	-9	-8	-7	-6	-5	-4	-3	-2	-1
T	h	i	s		i	s		a		s	e	n	t	e	n	c	e	.
0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18

- The length of a string is the number of characters it contains. The `len()` function returns the length of a string.
- We can extract individual characters by *indexing* the string at a given location. The first character in the string is located at index zero. If the length of the string is N, the final character is located at index N-1. Indexing outside string boundaries gives rise to an error.

```
s = 'This is a sentence.'
print(len(s))
print(s[0])
print(s[1])
print(s[2])
print(s[18])
print(s[19])
```

```
19
T
h
i
.
```

```
-----
IndexError                                Traceback (most recent call last)
Input In [4], in <module>
      5 print(s[2])
      6 print(s[18])
----> 7 print(s[19])

IndexError: string index out of range
```

- In Python it is possible to index relative to the end of the string using negative indices: The last character is at index -1, the second last at index -2, the third last at index -3, etc.

```
print(s[-1])
print(s[-2])
```

```
print(s[-3])
print(s[-19])
print(s[-20])
```

```
.
e
c
T
```

```
-----
IndexError                                Traceback (most recent call last)
Input In [5], in <module>
      3 print(s[-3])
      4 print(s[-19])
----> 5 print(s[-20])

IndexError: string index out of range
```

String slicing

- We can extract more than a single character from a string. We can extract subsequences or *slices* by specifying a range of indices separated by a colon. Writing `s[start:end]` will return a new string composed of the characters `s[start]`, `s[start+1]`, `s[start+2]`, ..., `s[end-3]`, `s[end-2]`, `s[end-1]`.
- Note that the character located at `s[end]` is *not* returned.

```
s = 'This is a sentence.'
print(s[0:4])
print(s[5:7])
```

```
This
is
```

- If either the starting or ending indices are omitted their values default to the beginning and end of the string.

```
s = 'This is a sentence.'
print(s[:4])
print(s[10:])
print(s[:])
```

```
This
sentence.
This is a sentence.
```

- As usual, negative indices can be used to specify locations relative to the end of the string.

```
s = 'This is a sentence.'  
print(s[:-10])
```

This is a

- Slicing is more forgiving in terms of indices than pure indexing. A start or end index outside of a string's boundaries is treated as the string boundary.

```
s = 'This is a sentence.'  
print(s[:100])  
print(s[-100:])  
print(s[-100:100])
```

This is a sentence.
This is a sentence.
This is a sentence.

Extended slicing

- It is possible to specify a third parameter when slicing sequences. It indicates the *step size* to take along the sequence when extracting its elements. Writing `s[start:end:step]` will return a new string composed of `s[start]`, `s[start+step]`, `s[start+2*step]`, `s[start+3*step]`, etc. Extraction continues for as long as `start+i*step < end` where `i = 0, 1, 2, 3`, etc.
- As usual, if the start or end of the slice are omitted they default to the beginning and end of the string respectively.

```
s = 'This is a sentence.'  
print(s[::1])  
print(s[::2])  
print(s[:10:3])  
print(s[10::3])
```

This is a sentence.
Ti sasnec.
Tss
stc

- A negative step size is interpreted as a step backwards through the sequence. This is handy for reversing a string.
- For a negative step size, if the start or end of the slice are omitted, their values default to the end and beginning of the string respectively.

```
s = 'This is a sentence.'
print(s[-1:-20:-1])
print(s[::-1])
print(s[::-2])
print(s[-18:-10:-1])
print(s[-10:-18:-1])
```

```
.ecnetnes a si sihT
.ecnetnes a si sihT
.censas iT

a si si
```

String concatenation and replication

- We can use the `+` operator to concatenate strings.

```
line1 = 'Humpty Dumpty sat on a wall'
line2 = 'Humpty Dumpty had a great fall'
print(line1 + ',\n' + line2 + '.')
```

```
Humpty Dumpty sat on a wall,
Humpty Dumpty had a great fall.
```

- We can use the `*` operator to replicate strings.

```
s = 'apple '
print(s * 3)
```

```
apple apple apple
```

String comparison

- Strings can be tested for equality with the `==` operator.

```
print('cat' == 'dog')
print('mouse' == 'mouse')
```

```
False
True
```

Strings are True or False

- The empty string `''` is interpreted as `False`.

```
s = ''  
if s:  
    print(True)  
else:  
    print(False)
```

False

- Any non-empty string is interpreted as `True`.

```
s = 'apple'  
if s:  
    print(True)  
else:  
    print(False)
```

True

Strings are iterable

- Because a string is an *iterable* sequence we can use a `for` loop to examine each of its characters in turn.

```
s = 'apple'  
for c in s:  
    print(c)
```

a
p
p
l
e

Strings are immutable

- Strings are *immutable*. This means they cannot be modified. If we try to modify a string we get an error.

```
s = 'apple'  
s[0] = 'A'  
print(s)
```

```
-----  
TypeError                                Traceback (most recent call last)  
Input In [18], in <module>  
      1 s = 'apple'  
----> 2 s[0] = 'A'  
      3 print(s)  
  
TypeError: 'str' object does not support item assignment
```

- If we want a “modify” a string we have to create a new one from the original.

```
s = 'apple'  
s = 'A' + s[1:]  
print(s)
```

Apple

Processing lines of text

- A common operation is to read in a line of text from a file, strip any surrounding whitespace and split the line into its constituent tokens.

```
line = 'This is a line of text'  
tokens = line.strip().split()  
print(tokens)
```

['This', 'is', 'a', 'line', 'of', 'text']

- By default, the `split()` method splits strings on whitespace but we can ask it to split on other characters.

```
line = 'This,is,a,line,of,text'  
tokens = line.strip().split(',')  
print(tokens)
```

['This', 'is', 'a', 'line', 'of', 'text']

- We can use the `join()` method to glue lists of words back together into a single string. We need to tell `join()` which glue character to use.

```
line = 'This,is,a,line,of,text'
tokens = line.strip().split(',')
print(tokens)
newline = ' '.join(tokens)
print(newline)
newline = '-'.join(tokens)
print(newline)
```

```
['This', 'is', 'a', 'line', 'of', 'text']
This is a line of text
This-is-a-line-of-text
```

String membership

- We can use the `in` operator to check whether one string is a substring of another.

```
s = 'This is a sentence.'
print('This' in s)
print('x' in s)
print('.') in s)
```

```
True
False
True
```

String methods

- Python comes with built-in support for a large set of common string operations. These operations are called `methods` and they define the things we can do with strings. We have looked only at a small number of Python's string methods.
- Calling `help(str)` in the Python shell or `pydoc str` in a Linux terminal outputs a list of these methods. We see the methods we can invoke on a string `s` include `capitalize()` (returns a capitalized version of `s`), `isdecimal()` (returns `True` if `s` contains only decimal characters), `lower()` (returns a new copy of `s` with all characters converted to lowercase), etc.
- Many editors and IDEs support autocompletion, whereby, if `s` is a string, placing your cursor after the dot in `s.` and hitting tab will provide you with a list of the string methods that can be invoked on `s`. Then simply select the one you want.
- Note that, because strings are immutable, calling a method on a string will *not* alter the string itself.

```
s = 'apple'
print(s.capitalize())
print(s)
```

```
Apple
apple
```


- Whenever you find it necessary to carry out some string processing, first look up the list of built-in string methods. There may be one that will help you with your task. There is no point writing code that duplicates what a built-in string method can do for you already.

```
help(str)
```

Help on class str in module builtins:

```
class str(object)
|   str(object='') -> str
|   str(bytes_or_buffer[, encoding[, errors]]) -> str
|
|   Create a new string object from the given object. If encoding or
|   errors is specified, then the object must expose a data buffer
|   that will be decoded using the given encoding and error handler.
|   Otherwise, returns the result of object.__str__() (if defined)
|   or repr(object).
|   encoding defaults to sys.getdefaultencoding().
|   errors defaults to 'strict'.
|
|   Methods defined here:
|
|   __add__(self, value, /)
|       Return self+value.
|
|   __contains__(self, key, /)
|       Return key in self.
|
|   __eq__(self, value, /)
|       Return self==value.
|
|   __format__(self, format_spec, /)
|       Return a formatted version of the string as described by format_spec.
|
|   __ge__(self, value, /)
|       Return self>=value.
|
|   __getattr__(self, name, /)
|       Return getattr(self, name).
|
|   __getitem__(self, key, /)
|       Return self[key].
|
|   __getnewargs__(...)
|
|   __gt__(self, value, /)
|       Return self>value.
|
|   __hash__(self, /)
|       Return hash(self).
|
|   __iter__(self, /)
|       Implement iter(self).
|
|   __le__(self, value, /)
|       Return self<=value.
|
|   __len__(self, /)
|       Return len(self).
|
|   __lt__(self, value, /)
|       Return self<value.
```

```

__mod__(self, value, /)
    Return self%value.

__mul__(self, value, /)
    Return self*value.

__ne__(self, value, /)
    Return self!=value.

__repr__(self, /)
    Return repr(self).

__rmod__(self, value, /)
    Return value%self.

__rmul__(self, value, /)
    Return value*self.

__sizeof__(self, /)
    Return the size of the string in memory, in bytes.

__str__(self, /)
    Return str(self).

capitalize(self, /)
    Return a capitalized version of the string.

    More specifically, make the first character have upper case and the rest lower
    case.

casefold(self, /)
    Return a version of the string suitable for caseless comparisons.

center(self, width, fillchar=' ', /)
    Return a centered string of length width.

    Padding is done using the specified fill character (default is a space).

count(...)
    S.count(sub[, start[, end]]) -> int

    Return the number of non-overlapping occurrences of substring sub in
    string S[start:end]. Optional arguments start and end are
    interpreted as in slice notation.

encode(self, /, encoding='utf-8', errors='strict')
    Encode the string using the codec registered for encoding.

    encoding
        The encoding in which to encode the string.
    errors
        The error handling scheme to use for encoding errors.
        The default is 'strict' meaning that encoding errors raise a
        UnicodeEncodeError. Other possible values are 'ignore', 'replace' and
        'xmlcharrefreplace' as well as any other name registered with
        codecs.register_error that can handle UnicodeEncodeErrors.

endswith(...)
    S.endswith(suffix[, start[, end]]) -> bool

    Return True if S ends with the specified suffix, False otherwise.
    With optional start, test S beginning at that position.
    With optional end, stop comparing S at that position.
    suffix can also be a tuple of strings to try.

expandtabs(self, /, tabsize=8)

```

Return a copy where all tab characters are expanded using spaces.

If tabsize is not given, a tab size of 8 characters is assumed.

`find(...)`

`S.find(sub[, start[, end]]) -> int`

Return the lowest index in `S` where substring `sub` is found, such that `sub` is contained within `S[start:end]`. Optional arguments `start` and `end` are interpreted as in slice notation.

Return -1 on failure.

`format(...)`

`S.format(*args, **kwargs) -> str`

Return a formatted version of `S`, using substitutions from `args` and `kwargs`. The substitutions are identified by braces ('{' and '}').

`format_map(...)`

`S.format_map(mapping) -> str`

Return a formatted version of `S`, using substitutions from `mapping`. The substitutions are identified by braces ('{' and '}').

`index(...)`

`S.index(sub[, start[, end]]) -> int`

Return the lowest index in `S` where substring `sub` is found, such that `sub` is contained within `S[start:end]`. Optional arguments `start` and `end` are interpreted as in slice notation.

Raises `ValueError` when the substring is not found.

`isalnum(self, /)`

Return True if the string is an alpha-numeric string, False otherwise.

A string is alpha-numeric if all characters in the string are alpha-numeric and there is at least one character in the string.

`isalpha(self, /)`

Return True if the string is an alphabetic string, False otherwise.

A string is alphabetic if all characters in the string are alphabetic and there is at least one character in the string.

`isascii(self, /)`

Return True if all characters in the string are ASCII, False otherwise.

ASCII characters have code points in the range U+0000-U+007F. Empty string is ASCII too.

`isdecimal(self, /)`

Return True if the string is a decimal string, False otherwise.

A string is a decimal string if all characters in the string are decimal and there is at least one character in the string.

`isdigit(self, /)`

Return True if the string is a digit string, False otherwise.

A string is a digit string if all characters in the string are digits and there is at least one character in the string.

`isidentifier(self, /)`

Return True if the string is a valid Python identifier, False otherwise.

Call `keyword.iskeyword(s)` to test whether string `s` is a reserved identifier, such as `"def"` or `"class"`.

`islower(self, /)`

Return True if the string is a lowercase string, False otherwise.

A string is lowercase if all cased characters in the string are lowercase and there is at least one cased character in the string.

`isnumeric(self, /)`

Return True if the string is a numeric string, False otherwise.

A string is numeric if all characters in the string are numeric and there is at least one character in the string.

`isprintable(self, /)`

Return True if the string is printable, False otherwise.

A string is printable if all of its characters are considered printable in `repr()` or if it is empty.

`isspace(self, /)`

Return True if the string is a whitespace string, False otherwise.

A string is whitespace if all characters in the string are whitespace and there is at least one character in the string.

`istitle(self, /)`

Return True if the string is a title-cased string, False otherwise.

In a title-cased string, upper- and title-case characters may only follow uncased characters and lowercase characters only cased ones.

`isupper(self, /)`

Return True if the string is an uppercase string, False otherwise.

A string is uppercase if all cased characters in the string are uppercase and there is at least one cased character in the string.

`join(self, iterable, /)`

Concatenate any number of strings.

The string whose method is called is inserted in between each given string. The result is returned as a new string.

Example: `'.'.join(['ab', 'pq', 'rs']) -> 'ab.pq.rs'`

`ljust(self, width, fillchar=' ', /)`

Return a left-justified string of length `width`.

Padding is done using the specified fill character (default is a space).

`lower(self, /)`

Return a copy of the string converted to lowercase.

`lstrip(self, chars=None, /)`

Return a copy of the string with leading whitespace removed.

If `chars` is given and not None, remove characters in `chars` instead.

`partition(self, sep, /)`

Partition the string into three parts using the given separator.

This will search for the separator in the string. If the separator is found, returns a 3-tuple containing the part before the separator, the separator itself, and the part after it.

If the separator is not found, returns a 3-tuple containing the original string and two empty strings.

`replace(self, old, new, count=-1, /)`

Return a copy with all occurrences of substring `old` replaced by `new`.

`count`

Maximum number of occurrences to replace.

-1 (the default value) means replace all occurrences.

If the optional argument `count` is given, only the first `count` occurrences are replaced.

`rfind(...)`

`S.rfind(sub[, start[, end]]) -> int`

Return the highest index in `S` where substring `sub` is found, such that `sub` is contained within `S[start:end]`. Optional arguments `start` and `end` are interpreted as in slice notation.

Return -1 on failure.

`rindex(...)`

`S.rindex(sub[, start[, end]]) -> int`

Return the highest index in `S` where substring `sub` is found, such that `sub` is contained within `S[start:end]`. Optional arguments `start` and `end` are interpreted as in slice notation.

Raises `ValueError` when the substring is not found.

`rjust(self, width, fillchar=' ', /)`

Return a right-justified string of length `width`.

Padding is done using the specified fill character (default is a space).

`rpartition(self, sep, /)`

Partition the string into three parts using the given separator.

This will search for the separator in the string, starting at the end. If the separator is found, returns a 3-tuple containing the part before the separator, the separator itself, and the part after it.

If the separator is not found, returns a 3-tuple containing two empty strings and the original string.

`rsplit(self, /, sep=None, maxsplit=-1)`

Return a list of the words in the string, using `sep` as the delimiter string.

`sep`

The delimiter according which to split the string.

None (the default value) means split according to any whitespace, and discard empty strings from the result.

`maxsplit`

Maximum number of splits to do.

-1 (the default value) means no limit.

Splits are done starting at the end of the string and working to the front.

`rstrip(self, chars=None, /)`

Return a copy of the string with trailing whitespace removed.

If `chars` is given and not None, remove characters in `chars` instead.

`split(self, /, sep=None, maxsplit=-1)`

Return a list of the words in the string, using `sep` as the delimiter string.

```

sep
    The delimiter according which to split the string.
    None (the default value) means split according to any whitespace,
    and discard empty strings from the result.
maxsplit
    Maximum number of splits to do.
    -1 (the default value) means no limit.

splitlines(self, /, keepends=False)
    Return a list of the lines in the string, breaking at line boundaries.

    Line breaks are not included in the resulting list unless keepends is given as
    true.

startswith(...)
    S.startswith(prefix[, start[, end]]) -> bool

    Return True if S starts with the specified prefix, False otherwise.
    With optional start, test S beginning at that position.
    With optional end, stop comparing S at that position.
    prefix can also be a tuple of strings to try.

strip(self, chars=None, /)
    Return a copy of the string with leading and trailing whitespace removed.

    If chars is given and not None, remove characters in chars instead.

swapcase(self, /)
    Convert uppercase characters to lowercase and lowercase characters to uppercase.

title(self, /)
    Return a version of the string where each word is titlecased.

    More specifically, words start with uppercased characters and all remaining
    cased characters have lower case.

translate(self, table, /)
    Replace each character in the string using the given translation table.

    table
        Translation table, which must be a mapping of Unicode ordinals to
        Unicode ordinals, strings, or None.

    The table must implement lookup/indexing via __getitem__, for instance a
    dictionary or list. If this operation raises LookupError, the character is
    left untouched. Characters mapped to None are deleted.

upper(self, /)
    Return a copy of the string converted to uppercase.

zfill(self, width, /)
    Pad a numeric string with zeros on the left, to fill a field of the given width.

    The string is never truncated.

```

Static methods defined here:

```

__new__(*args, **kwargs) from builtins.type
    Create and return a new object. See help(type) for accurate signature.

maketrans(...)
    Return a translation table usable for str.translate().

    If there is only one argument, it must be a dictionary mapping Unicode
    ordinals (integers) or characters to Unicode ordinals, strings or None.
    Character keys will be then converted to ordinals.

```

| If there are two arguments, they must be strings of equal length, and
| in the resulting dictionary, each character in x will be mapped to the
| character at the same position in y. If there is a third argument, it
| must be a string, whose characters will be mapped to None in the result.